

# Kubernetes




A short introduction

# Before We Begin



## Requirements:

- **Minikube and or K3S**  
<https://github.com/kubernetes/minikube>
- **Virtualbox\*:**  
<https://www.virtualbox.org/wiki/Downloads>
- **kubectl:**  
<https://kubernetes.io/docs/tasks/tools/install-kubectl/>
- **k8s-intro-tutorials repo:**  
<https://github.com/mrbobbytables/k8s-intro-tutorials>



# Acknowledgements



These slides are an updated/modified version of the following ones:

[Introduction to Kubernetes - Presentazioni Google](#) from Bob Killen

We (SC/ Marco Cesaratto/ Ruggero Lot) updated and modified it for the course Advanced Cloud computing.

We thanks Bob for his great work and to allow modifications and updates under CC by 4.0

# Before We Begin



**Go here for all info related to the course**



[Foundations-of-HPC/Cloud-advanced-2023](#)  
[\(github.com\)](#)

# Kubernetes



A short introduction




# Project Overview

# What Does “Kubernetes” Mean?



Greek for “pilot” or  
“Helmsman of a ship”



[Image Source](#)





# What is Kubernetes?

- Project that was spun out of Google as an **open source container orchestration platform**.
- Built from the lessons learned in the experiences of developing and running Google's Borg and Omega.
- Designed from the ground-up as a **loosely coupled** collection of components centered around deploying, maintaining and scaling workloads.





# What Does Kubernetes do?

- Known as the **linux kernel of distributed systems**.
- **Abstracts away the underlying hardware** of the nodes and provides a uniform interface for workloads to be both deployed and consume the shared pool of resources.
- Works as an engine for resolving state by converging actual and the **desired state** of the system.

# Decouples Infrastructure and Scaling



- **All services** within Kubernetes are natively Load Balanced.
- Can scale up and down dynamically.
- Used both to enable self-healing and seamless upgrading or rollback of applications.

# Self Healing



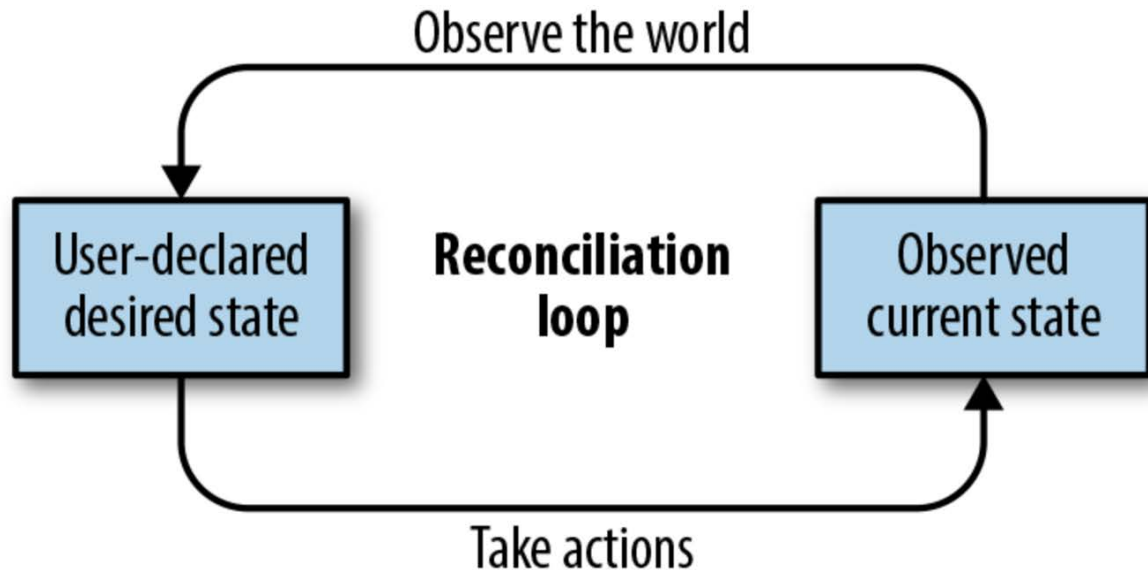
Kubernetes will **ALWAYS** try and steer the cluster to its desired state.

- **Me:** “I want 3 healthy instances of redis to always be running.”
- **Kubernetes:** “Okay, I’ll ensure there are always 3 instances up and running.”
- **Kubernetes:** “Oh look, one has died. I’m going to attempt to spin up a new one.”

# Self Healing



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



# What can Kubernetes REALLY do?



- Autoscale Workloads
- Blue/Green Deployments
- Fire off jobs and scheduled cronjobs
- Manage Stateless and Stateful Applications
- Provide native methods of service discovery
- Easily integrate and support 3rd party apps

# Most Importantly...



Use the **SAME** API  
across bare metal and  
**EVERY** cloud provider!!!

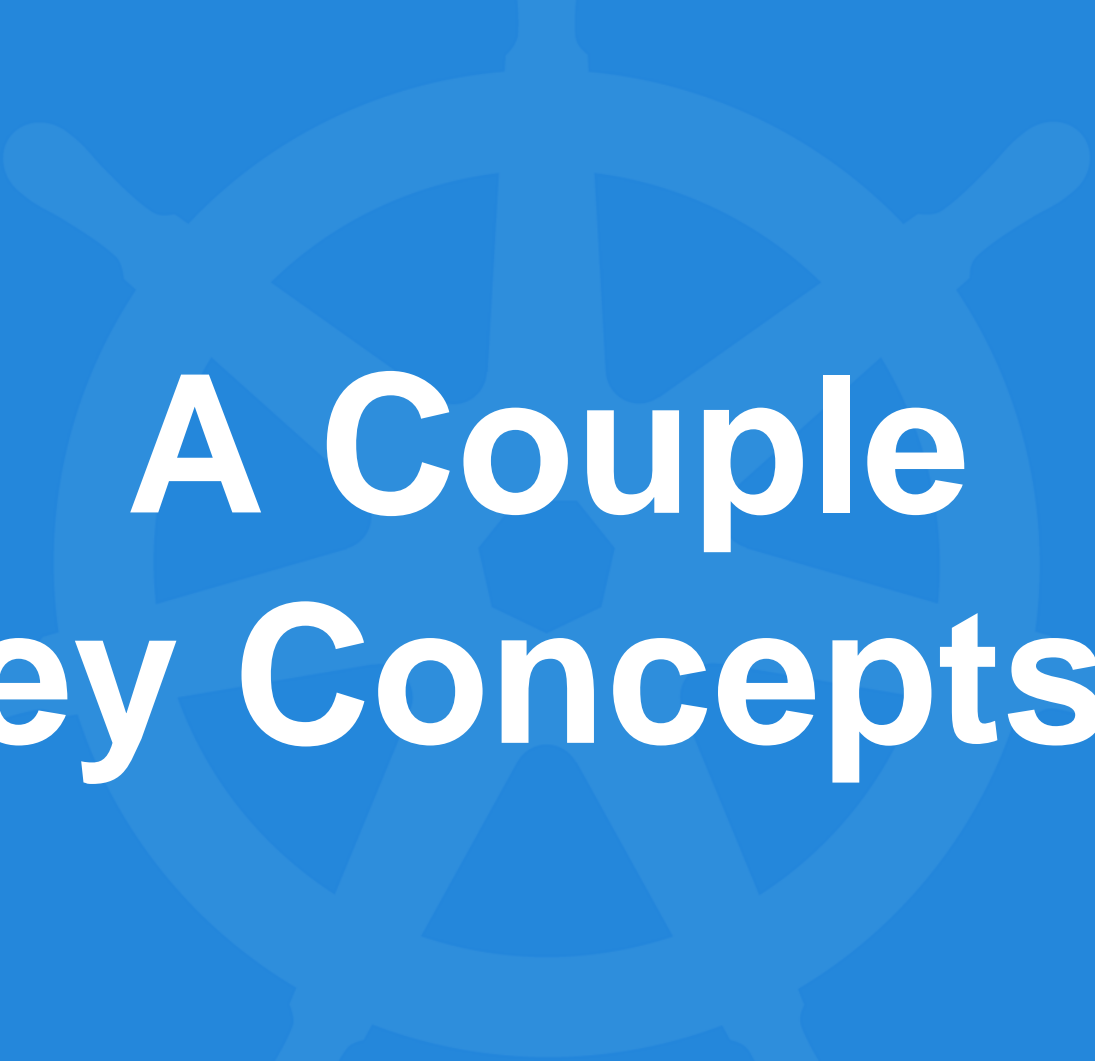


# Who “Manages” Kubernetes?



**CLOUD NATIVE**  
**COMPUTING FOUNDATION**

The CNCF is a child entity of the Linux Foundation and operates as a vendor neutral governance group.




# A Couple Key Concepts...



# Pods




- **Atomic unit** or smallest “*unit of work*” of Kubernetes.
- Pods are **one or MORE containers** that share volumes, a network namespace, and are a part of a **single context**.



# Pods



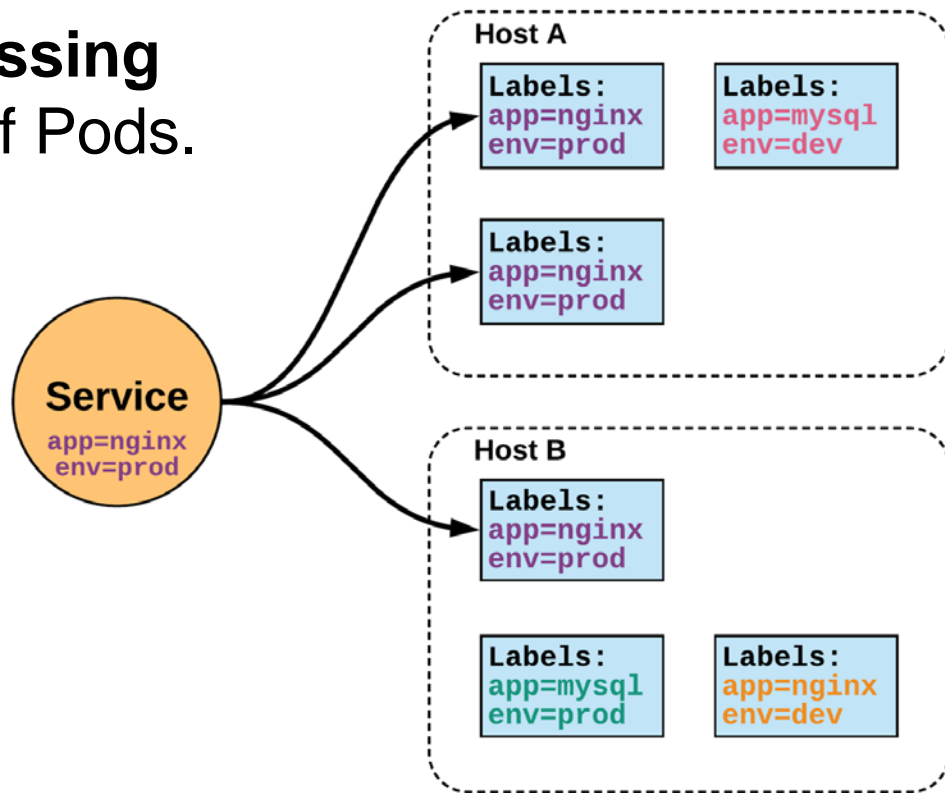
**They are  
also  
Ephemeral!**



# Services



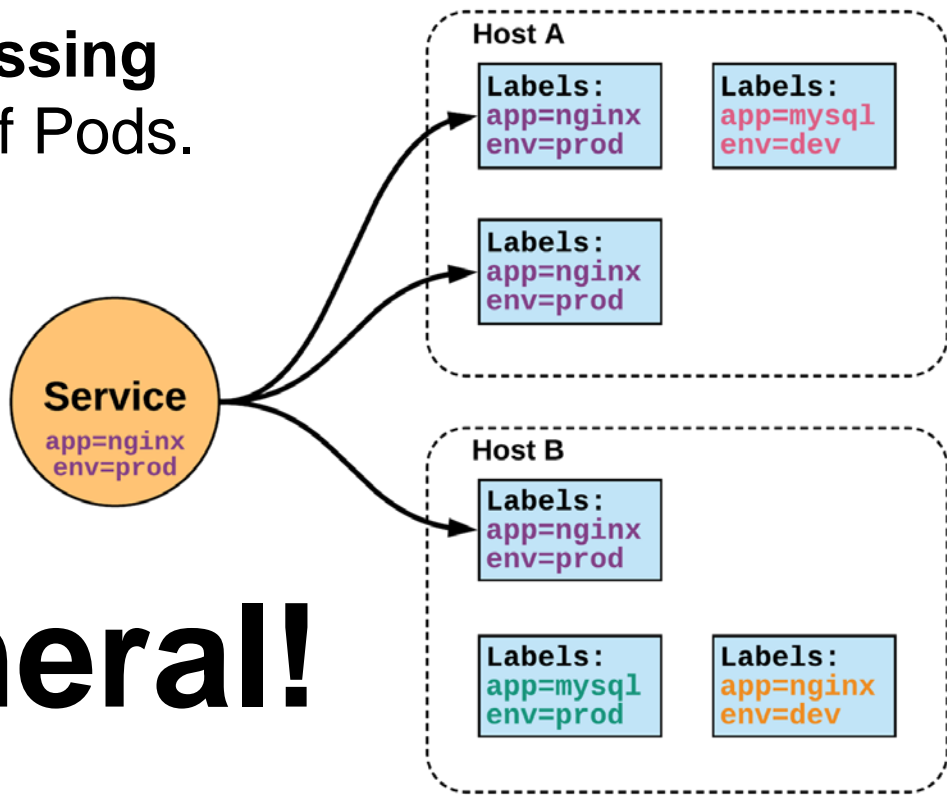
- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource**
  - static cluster IP
  - static namespaced DNS name



# Services



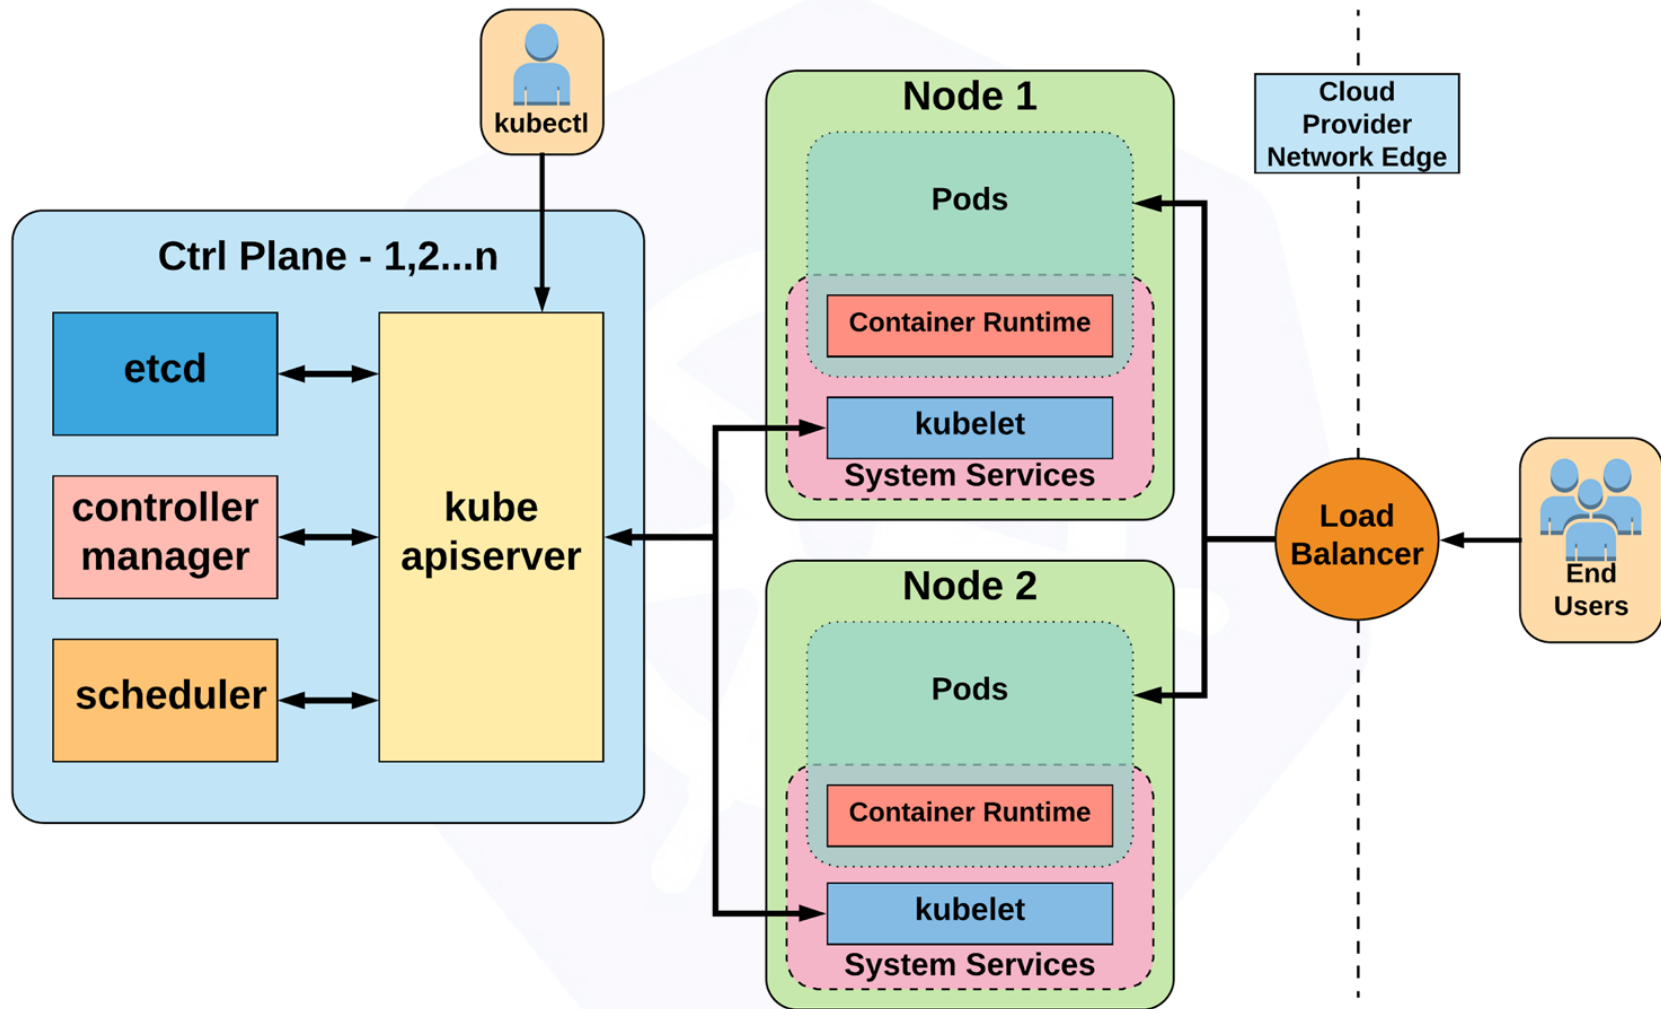
- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource**
  - static cluster IP
  - static namespaced DNS name



# NOT Ephemeral!



# Architecture Overview





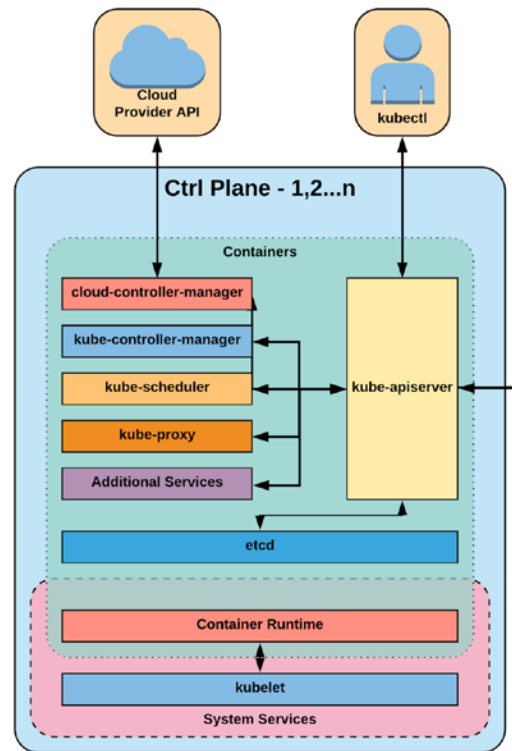
# Control Plane Components

## Architecture Overview

# Control Plane Components



- kube-apiserver
- etcd
- kube-controller-manager
- kube-scheduler





# kube-apiserver



- Provides a forward facing REST interface into the kubernetes control plane and datastore.
- All clients and other applications interact with kubernetes **strictly** through the API Server.
- Acts as the gatekeeper to the cluster by handling authentication and authorization, request validation, mutation, and admission control in addition to being the front-end to the backing datastore.

# etcd




- etcd acts as the cluster datastore.
- Purpose in relation to Kubernetes is to provide a strong, consistent and highly available key-value store for persisting cluster state.
- Stores objects and config information.





Uses “*Raft Consensus*” among a quorum of systems to create a fault-tolerant consistent “*view*” of the cluster.

<https://raft.github.io/>





# kube-controller-manager


- Serves as the primary daemon that manages all core component control loops.
- Monitors the cluster state via the apiserver and **steers the cluster towards the desired state.**

List of core controllers: <https://github.com/kubernetes/kubernetes/blob/master/cmd/kube-controller-manager/app/controllermanager.go#L344>

# kube-scheduler



- Verbose policy-rich engine that evaluates workload requirements and attempts to place it on a matching resource.
- Default scheduler uses bin packing.
- Workload Requirements can include: general hardware requirements, affinity/anti-affinity, labels, and other various custom resource requirements.




# **Node Components**

Architecture Overview

# Node Components



- kubelet
- kube-proxy
- Container Runtime Engine



# kubelet



- Acts as the node agent responsible for managing the lifecycle of every pod on its host.
- Kubelet understands YAML container manifests that it can read from several sources:
  - file path
  - HTTP Endpoint
  - etcd watch acting on any changes
  - HTTP Server mode accepting container manifests over a simple API.



# kube-proxy

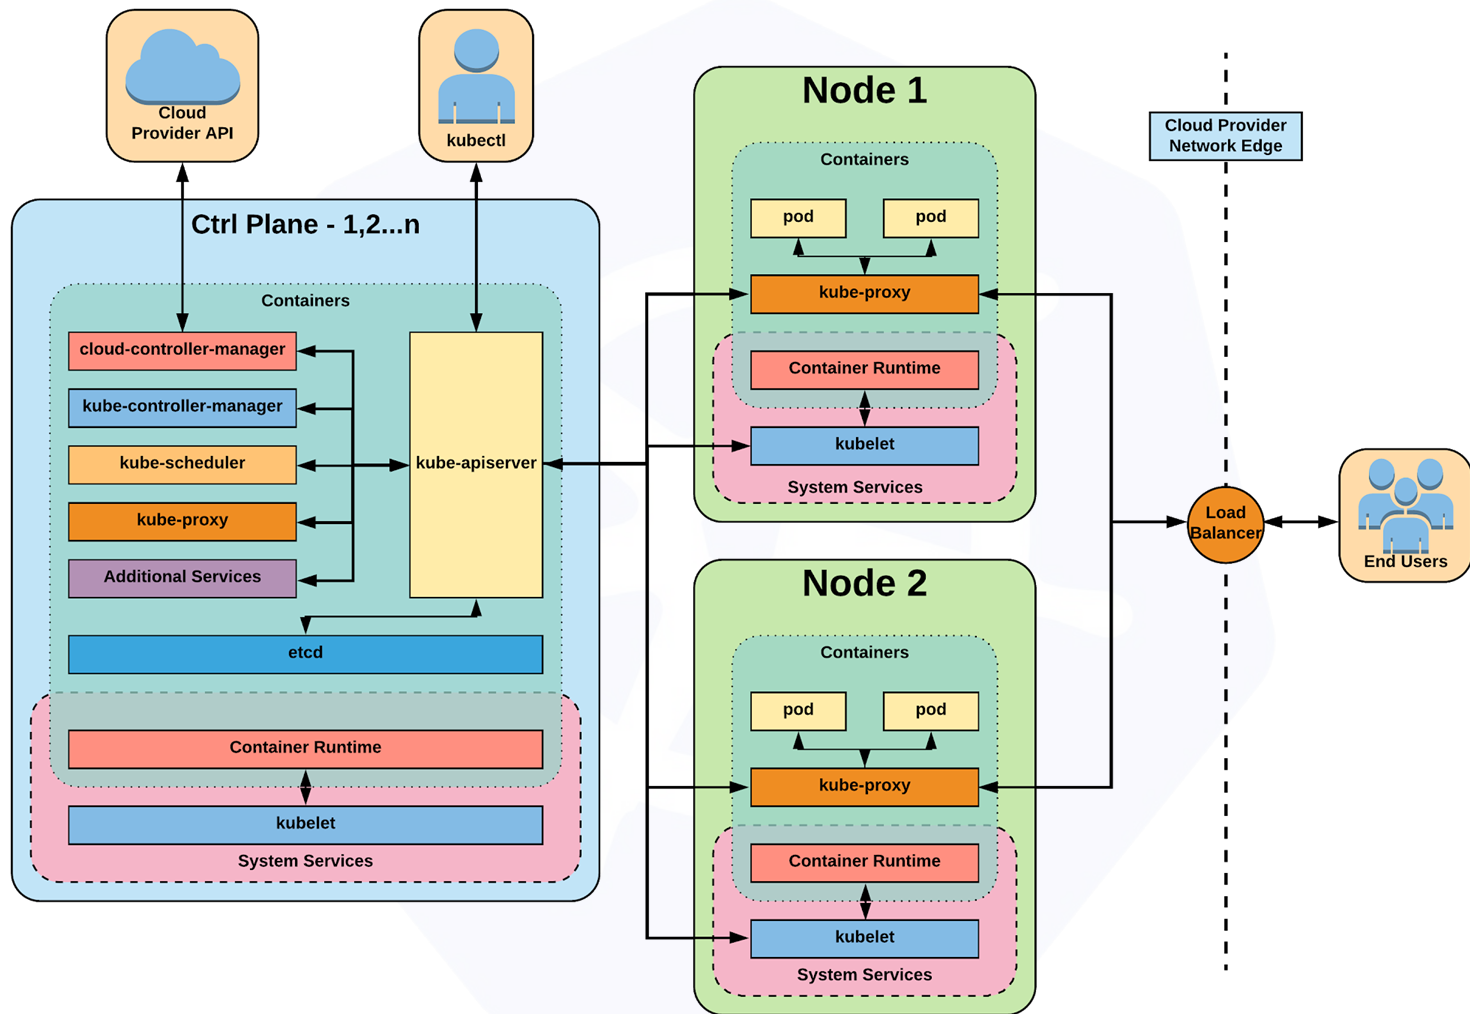


- Manages the network rules on each node.
- Performs connection forwarding or load balancing for Kubernetes cluster services.
- Available Proxy Modes:
  - Userspace
  - iptables
  - ipvs (default if supported)

# Container Runtime Engine



- A container runtime is a CRI (Container Runtime Interface) compatible application that executes and manages containers.
  - Containerd (docker)
  - Cri-o
  - Rkt
  - Kata (formerly clear and hyper)
  - Virtlet (VM CRI compatible runtime)





**End of the first part**

# Links



- **Free Kubernetes Courses**

<https://www.edx.org/>

- **Interactive Kubernetes Tutorials**

<https://www.katacoda.com/courses/kubernetes>

- **Learn Kubernetes the Hard Way**

<https://github.com/kelseyhightower/kubernetes-the-hard-way>

- **Official Kubernetes Youtube Channel**

<https://www.youtube.com/c/KubernetesCommunity>

- **Official CNCF Youtube Channel**


<https://www.youtube.com/c/cloudnativefdn>

- **Track to becoming a CKA/CKAD (Certified Kubernetes Administrator/Application Developer)**

<https://www.cncf.io/certification/expert/>

- **Awesome Kubernetes**

<https://ramitsurana.gitbooks.io/awesome-kubernetes/content/>





# Questions?