Action plan: Installing Kubernetes

↝ Some considerations

→ •) Access by many user ↝ we need some control magent
•) Consistetly available ↝
•) adaptable / Scalable on Demand.
•) HA ⇒ Distributed {in every Part}

↝ Architecture

More?? on
kube Proxy

the parts of a kubernets cluster

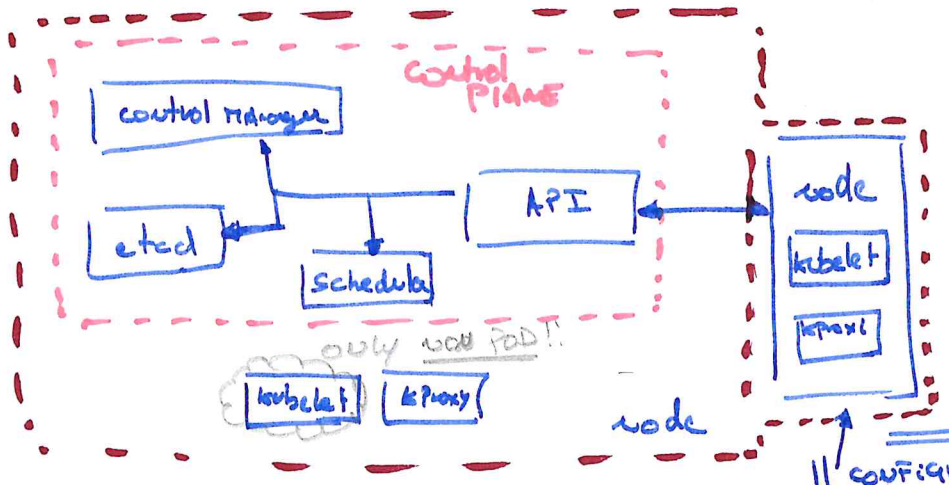

GOAL!!

Simpler installation:
everythig on the same node

the PODS but not only PODS

|| configured to run kubernetes Pods ||

↝ resources in kubernetes are again organized in namespaces. ≠ kind namespace

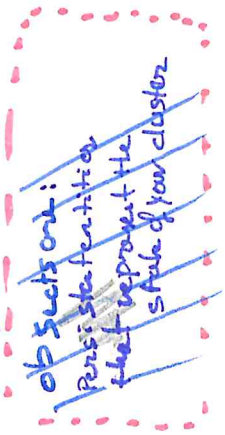

kube-System namespace } → base namespace where it is possible to find All the managin resources } → possible to

control-manager : Loop that watches the states of the kluster (by looking into the API Server) And move it towards the Desired State

API-server : Validate new Data (object) in input, Provides a front-end for the cluster shared state

kube-proxy : kubernetes network proxy

Scheduler : Assign PoD to nodes. The scheduler determines which nodes are valid placements accordingly to available resources.

kubelet: Responsible for registering the node to the API Server

→ One of the few non-container applications PART of kubernetes

note: All these pods (Special pods) share the same NetNS or pid ns

=> you can check it by observing that

- kube-proxy
- kube-API
- kube-scheduler
- etcd

are opening ~~ports~~ [network] ports on the node!!

[ss -tunalp]

Bird
CAlico ?
CEPH

Starting from our VM we are going to **Install Kubernetes**
↳ 2core + 2GB ram

This first installation will **not be HA**, later on we will see How to move from a basic installation to HA.

Prefered Method ⇒ **kube ADM**

But what is kubernetes

- K8S → by GOOGLE (....)
- K3S → by RANCHER (SUSE)
  (Light kubernetes
  less components, eg
  Etcd → Sqlite)
- K0S → even smaller than K3S
  (young project)

.Prepare the field for **cri-o** & **kubelet**

kernel modules

- ~> overlay → enable overlay fs kernel module
- ~> br_net_filter → Allow Net Filter @ Bridge Level.

kernel Parameters

- bridge_nf_call_* ⇒ enable nf tables calls (for Firewall @ bridge level)
- ip4.ip_forward ⇒ Allows **pkgs forwarding**
  ↳ GET PKGS in from an interface and throw them out of another interface.

kernel Parameters.

Cgroups: they perform **resource control** and are **managed by SYSTEMD**
↳ {How much ram / % CPU a Process can use?} SAME as **SLURM**

later

Miscellaneous

Disable: SeLinux → 2 security modules on top of everything we have seen
- SeLinux
- App Armor

} 2 diffent philosophies, kub only support Apparmor for now

Disable t-ram → a kind of swap in ram.

CRI-O configuration

/etc/crio/ → Pure crio conf. { • CNI is coded here partially

/etc/containers/registries.conf for containers (reg. stuff, storage conf, ...)

/etc/cni/net.d/* ← real CNI configuration!!

type: bridge
bridge: cni0
...
more...

} Simplest Plugin.

obs on worker nodes?
obs on user Management?
Limits on worker nodes?

Limitation of this plugin: no inter node communication between pods

Very Good Starting point anyway!

FINAL INSTALLATION

KUBEADM

-- pod network -cidr
-- services - cidr

Classes Inter Domain Routing

Just refer to the Stadard notation

* we need some more knowledge to set this up... Calico...

-- control-plane
-endpoint

the ips of our pods! :-) Must be = to what we declared to cni! CNI!

we need to tell kub. where is the control plane!

this should point to a DNS name or better to an External Load balancer over All the control planes we are hosting.

control plane HA

Hurray!! our "cluster" is up & running!!

Show that cnio- is playing a central role in cluster operations.

To the point that we can bypass kubelet to ask the status of the SINGLE node

kubectl is the tool to explore the cluster we just created:

kube ctl version : dos: ±3 from kubernetes API minor version: 1.1.0 (minor)

states of the cluster

kube ctl get --raw = /healthz [ ] → if
= /livez
= /readyz
optional. ? verbose

it is better to define what they are used for

if FAILS the container is killed And restarted by kubelet

SAME AS HEALTH but Act BEFORE

when is ok (200) your application starts to receive traffic

Where is this information used?

well written application should expose All these 3 endpoints in their API in order to show

~> Livez => that they're booting has ended correctly
Probed by KUBECTL as Startup Probe
Longer failure threshold Between diffrt checks Are usually Done [Act As liveness]

~> readyz => that they are ready to accept traffic
readiness Prob once good, your APP will start receive traffic from outside

~> Healthz => that the app is working as expected
Liveness Probe
When it fails the container is killed Abruptly and restarted.

RANDOM NAMES
can be API endpoints or files created at runtime, kubernetes Doesn't care as long as it is instructed to probe them.

Defined in the object as part of the SPEC of CONTAINER